

NEXUS - Master Engineering Thesis and Research Program

INDRANIL

26 September 2026

Contents

Executive compression	1
1. Document status and claim discipline	2
1.1 Why this document exists	2
1.2 The current maturity boundary	2
2. The central thesis	3
2.1 From screen-centric computing to intent-centric computing	3
2.2 The new unit of computing	3
2.3 The system-level objective	3
3. Philosophical premise: the machine should meet the human where the human is	4
4. Nexus as a dynamical system	4
5. The semantic operating system	5
5.1 Core objects	5
World	5
Context	5
Memory	6
Intent	6
Capability	6
Action	6
Policy	6
Compute	6
Trust	6
5.2 Intent IR	6
5.3 Action IR	7
5.4 Capability graph	7
5.5 Compatibility is a ladder, not a rewrite	7
6. Memory as a longitudinal personal substrate	8
6.1 Memory types	8
6.2 Memory write policy	8
7. Human signal architecture	9
7.1 Multimodality is a design principle, not a feature list	9
7.2 Gaze	9
7.3 EMG	10
7.4 EEG	10
7.5 Voice, vision, and IMU	10

8. Neuroengineering and bio-nano architecture	10
8.1 The hierarchy of biological interfaces	10
8.2 Electrode stack	11
8.3 Candidate materials	11
8.4 Future invasive path	11
9. Intelligence architecture	12
9.1 The assistant is a system process	12
9.2 From prompt-response to persistent policy	12
9.3 Predictive interaction	12
10. Continual learning and post-training	13
10.1 Three adaptation levels	13
Level 1: memory adaptation	13
Level 2: user-specific policy adaptation	13
Level 3: global model improvement	13
10.2 Safe post-training pipeline	13
10.3 Why online RL must be bounded	13
11. Recursive epistemic intelligence	14
11.1 Working definition	14
11.2 The recursion ladder	14
11.3 Fixed search space versus changing the space	14
11.4 Falsification behavior	15
12. World models and reality-facing intelligence	15
13. The compute reservoir	15
13.1 Why the reservoir exists	15
13.2 R0, R1, R2	16
R0 - pocket reservoir	16
R1 - portable workstation reservoir	16
R2 - stationary reservoir	16
13.3 Compute placement as an optimization problem	16
13.4 Wireless and wired modes	16
Class A - reflex	17
Class B - interaction	17
Class C - intelligence	17
13.5 Transport the minimum sufficient representation	17
14. Power and thermal engineering	17
15. Photonics: the physical interface as a computer	18
15.1 Optical architecture invariant	18
15.2 Field representation	18
15.3 Waveguide physics	18
15.4 Etendue and the unavoidable trade space	19
15.5 External evidence for the direction	19
16. Adaptive optical control	19
16.1 State-dependent optical policy	19
16.2 Optical co-design with intelligence	20
16.3 Computational foveation	20
17. Optical see-through and the 3L interface	20
17.1 Current product architecture choice	20

17.2 Nexus 3L	21
18. Eye, depth, and perceptual constraints	21
19. Optical safety architecture	21
20. The bioelectronic safety hierarchy	22
21. Security and neurodata sovereignty	22
21.1 Data classes	22
21.2 Threat model	22
21.3 Privacy as a physical architecture	23
22. Human effort as a measurable systems variable	23
23. OS runtime and developer platform	23
24. The capability ecosystem	24
25. The physical product architecture	25
25.1 Glasses	25
25.2 Wrist band	25
25.3 Optional audio device	25
25.4 Reservoir	25
26. Distributed execution and graceful degradation	25
27. The core control law of Nexus	26
28. Benchmark suite	26
28.1 Intelligence benchmark	26
28.2 Neural benchmark	27
28.3 Optical benchmark	27
Geometric	27
Radiometric	27
Photometric	27
Wave-optic	27
Temporal	27
Human	28
29. Decisive experiments	28
Experiment 1 - Does context outperform sensor complexity?	28
Experiment 2 - Does prediction reduce total human effort?	28
Experiment 3 - Does personalization beat model scaling?	28
Experiment 4 - Can intelligent optics reduce resource cost?	28
Experiment 5 - Does the reservoir improve the wearable system frontier?	29
Experiment 6 - Can the intelligence discover and reject wrong models?	29
30. Validation ladder	29
31. Falsifiability and failure modes	29
31.1 The optical thesis fails if	29
31.2 The neural thesis fails if	30
31.3 The intelligence thesis fails if	30
31.4 The reservoir thesis fails if	30
31.5 The platform thesis fails if	30

32. Evidence from current science and engineering	30
33. The architecture moat	31
34. Open versus proprietary boundaries	31
35. Manufacturing and systems engineering	31
36. The product evolution path	32
Stage A - semantic Nexus on commodity hardware	32
Stage B - multimodal wearable interaction	32
Stage C - reservoir-native execution	32
Stage D - advanced optical co-design	32
Stage E - neural expansion	32
Stage F - longitudinal intelligence	32
Stage G - advanced neurotechnology research	32
37. The research operating model	32
38. The engineering creed	33
39. Master system optimization	33
40. The single deepest bottleneck	34
41. Why the architecture can still be radical without pretending everything is solved	34
42. Public research agenda	35
Software	35
Intelligence	35
Neurotechnology	35
Photonics	35
Systems	35
43. What Nexus is, finally	36
44. Final compression	36
45. A researcher's reading of Nexus	37
46. Cognitive runtime, attention, and externalized working memory	37
46.1 Externalized working memory	38
46.2 Attention must remain user-steerable	38
47. Applications as experimental environments	39
47.1 Scientific research	39
47.2 CAD and engineering	39
47.3 Education	39
47.4 Communication	39
47.5 Industrial and field operations	39
48. What is deliberately not the novelty claim	40
49. Current Nexus branch decisions versus architectural invariants	40
Current product branch	40
Architectural invariants	40
50. Master evidence matrix	41

Appendix N. Formal state machine for a Nexus action	41
Appendix O. Failure taxonomy	42
F0 - sensor failure	42
F1 - estimation failure	42
F2 - context failure	42
F3 - model failure	42
F4 - policy failure	42
F5 - capability failure	43
F6 - transport failure	43
F7 - actuation failure	43
F8 - safety failure	43
Appendix P. Resource-aware inference policy	43
Appendix Q. Research log schema	43
Appendix R. Final public disclosure	44
Appendix A. Minimal Intent IR schema	44
Appendix B. Minimal Action IR schema	44
Appendix C. Protocol and compute classes	45
Appendix D. Neurotechnology benchmark protocol	45
Appendix E. Optical parameterization	45
Appendix F. Reservoir thermal model	46
Appendix G. Current public evidence sources	46
Appendix H. Nexus internal source corpus integrated into this edition	47
Appendix I. Master architecture figure	48
Appendix J. Three-timescale figure	49
Appendix K. Photonic architecture figure	50
Appendix L. Learning loop figure	51
Appendix M. Evidence ladder figure	52

NEXUS

Master Engineering Thesis and Research Program

A Distributed Human-Compute Operating Environment

From multimodal human intent to adaptive computation and programmable photonic action

Physics | Photonics | Bioelectronics | Neurotechnology | AI | OS | Control | Safety

Public Technical Edition v1.0

26 September 2026

Research architecture / engineering program / public discussion document

This document is a research and engineering thesis, not a claim that the complete Nexus consumer system already exists. Established physical laws, published external demonstrations, current Nexus implementations, system hypotheses, and unresolved engineering targets are deliberately separated. The architecture is intended to be falsifiable: every hard claim should terminate in a measurement, benchmark, or failure criterion.

Executive compression

Nexus begins from a simple observation: the computer became extraordinarily capable while the human interface remained organized around explicit device operation. We still translate intention into app selection, menus, windows, cursors, touch targets, typing, and repeated attention switching. Artificial intelligence improves what computation can do, but it does not automatically remove the interaction architecture that sits between a human goal and a computational capability.

The Nexus thesis is therefore not “build better smart glasses.” It is:

Build a persistent human-compute operating environment in which human state, world state, memory, intent, capability, computation, and physical presentation are first-class system primitives.

The physical endpoint is advanced everyday eyewear, but the architecture is distributed. The glasses are the sensory and optical organ. A personal compute reservoir provides persistent local intelligence, memory, adaptation, and heavier inference without forcing the frame to carry the entire thermal and energy budget. Nearby wired systems and remote infrastructure extend compute when latency, privacy, bandwidth, or workload demand it.

The software layer is the semantic core:

```
human signals
  -> state estimation
  -> intent evidence
  -> Intent IR
  -> context + memory
  -> capability graph
  -> policy / plan
  -> Action IR
  -> compute allocation
  -> optical / audio / device actuation
  -> measured outcome
  -> memory + bounded learning
```

The deepest architectural shift is that the interface is not treated as a passive input/output shell. The optical system itself becomes a controllable physical subsystem. Human state, task state, world state, and available computation can jointly determine what should be rendered, where, at what detail, and with what optical resource. The research problem becomes a closed-loop optimization over information, optics, compute, power, thermal load, safety, and human effort.

The program has several hard edges. EEG is noisy and nonstationary. EMG is promising but still requires robust cross-session and cross-user modeling. Optical see-through waveguides must simultaneously satisfy field of view, eyebox, efficiency, distortion, brightness, color, uniformity, latency, safety, mass, and manufacturability. A reservoir must provide enough compute to make the system intelligent without turning the product into a belt of hot batteries and cables. Continual learning must become genuinely useful without becoming silently manipulative, privacy-invasive, or safety-unsafe. None of these problems disappears because the architecture is elegant.

That is precisely why the thesis is written as an engineering program rather than a product manifesto.

The invariant is not a particular laser, waveguide, chip, electrode, model, or wireless standard. The invariant is the loop:

human/world/task state → inference → capability → physical action → measured consequence → updated model

Everything else is replaceable until evidence says otherwise.

1. Document status and claim discipline

1.1 Why this document exists

The Nexus project has been developed across separate streams: a semantic operating system, a computing reservoir, an all-in-one product architecture, a photonics and optics research stack, a neural and physiological interface program, a personal intelligence architecture, a recursive epistemic intelligence theory, and a practical execution plan. The purpose of this master thesis is to collapse those streams into one technical object without flattening the important distinctions between them.

A public technical document has to do more than sound visionary. It must answer six questions continuously:

1. What is physically established?
2. What has been demonstrated externally by other researchers or companies?
3. What has actually been implemented inside Nexus?
4. What remains a Nexus hypothesis?
5. What experiment would distinguish the hypothesis from its competitors?
6. What happens if the result is negative?

The project therefore uses the following evidence language.

Label	Meaning	Public wording
P0	Physical law / standard mathematical result	established principle
P1	Published external demonstration	externally demonstrated
N1	Nexus software architecture / prototype	implemented in Nexus
N2	Nexus simulation / hardware model	simulated / modeled
H1	Research hypothesis	proposed hypothesis
R1	Required future result	engineering requirement
U1	Unknown / unresolved	open research problem

The strongest claims in this thesis are often the ones that remain explicitly marked H1 or U1. A project becomes less credible when uncertainty is hidden behind fluent language.

1.2 The current maturity boundary

At the date of this edition, Nexus has a meaningful software and research architecture. The project includes Intent IR and Action IR abstractions, context/world/memory models, compute-fabric abstractions, BCI-facing interfaces, optical-field simulation primitives, and a small closed-loop reality-facing experimental program using predictive latent representations and learned dynamics. Those artifacts demonstrate architectural intent and research capability; they do not demonstrate a finished consumer product.

The following remain research or engineering targets rather than completed product facts:

- final glasses mechanical design;
- production optical engine;
- production laser/PIC or alternate source architecture;
- prescription-ready optical packaging;
- robust long-duration dry EEG on consumer eyewear;
- final EMG and EEG sensor placement;
- universal neural decoding;
- production silicon;
- final battery and thermal envelope;
- large-scale longitudinal human validation;
- regulatory certification where applicable;
- manufacturing yield at consumer economics.

The correct posture is therefore neither false modesty nor false certainty:

conviction in the problem + humility about the unsolved mechanisms

2. The central thesis

2.1 From screen-centric computing to intent-centric computing

The old abstraction can be simplified as:

$$H \rightarrow P \rightarrow GUI \rightarrow OS \rightarrow App \rightarrow Compute$$

where the person must repeatedly translate a goal into device-specific operations.

Nexus proposes:

$$H \rightarrow Intent \rightarrow Capability \rightarrow Compute \rightarrow Action$$

The difference is not that keyboards, touchscreens, or conventional operating systems disappear. They remain capabilities inside a larger semantic environment. The change is that the human-facing operating model no longer requires the user to think in the same categories as the hardware or application packaging.

A person can say “compare these two papers,” look at two objects, make a small motor signal, or use a conventional keyboard. The application receives a stable semantic request rather than five separate input implementations.

2.2 The new unit of computing

The conventional unit is the application process. Nexus adds another unit:

Human task episode

A task episode contains:

- the user’s state and explicit goal;
- relevant world state;
- contextual history;
- memory references;
- candidate capabilities;
- compute resources;
- actions taken;
- outcomes;
- corrections;
- learned consequences.

This gives the operating layer a temporal dimension. The user does not restart at zero after closing an app.

2.3 The system-level objective

The primary product objective can be represented as:

$$J = U_{task} - \lambda_v C_{visual} - \lambda_c C_{cognitive} - \lambda_i C_{interaction} - \lambda_e C_{energy} - \lambda_r C_{risk}$$

subject to hard constraints on safety, thermal state, latency, privacy, and physical feasibility.

A useful operational metric is:

$$U_I = \frac{N_{\text{successful outcomes}}}{C_{\text{explicit interaction}}}$$

This is not a universal intelligence metric. It is a product-level measure of whether Nexus is removing meaningful interface work rather than merely adding more technology.

3. Philosophical premise: the machine should meet the human where the human is

Technology often forces the human to become an interpreter of technology. Nexus reverses that direction.

The system should estimate the user's state instead of requiring the user to translate every intention into a rigid sequence of controls. But this is not an argument for an omniscient machine. The system must remain corrigible, inspectable, and subordinate to explicit authority boundaries.

The philosophical compression is:

The user should spend attention on the world and the goal, not on operating the computer.

The Vedantic reading of the architecture is useful only as a metaphor, not as a scientific premise: the interface should subtract unnecessary layers between the observing person and the action being performed. Neti-neti becomes an engineering heuristic here: remove everything that is not necessary to accomplish the task, then measure whether the remaining system is actually better.

This creates an unusual design principle:

less visible interface $\neq$ less computation

The complexity moves into the system so that conscious complexity can move out of the human loop.

That movement is valuable only if the machine earns the right to act by being measurably reliable.

4. Nexus as a dynamical system

The master mathematical abstraction is a partially observed, constrained, multimodal dynamical system.

Let:

H_t = human state

W_t = world state

M_t = persistent memory state

Q_t = task and goal state

C_t = compute / energy / network state

O_t = optical and presentation state

A_t = authorized physical or digital action

The system receives observations:

$$Y_t = \{Y_t^{EEG}, Y_t^{EMG}, Y_t^{gaze}, Y_t^{voice}, Y_t^{vision}, Y_t^{IMU}, Y_t^{device}\}.$$

The latent belief state is:

$$b_t(x) = P(X_t = x \mid Y_{0:t}, A_{0:t-1}, M_t).$$

An action updates the physical state:

$$X_{t+1} \sim P(X_{t+1} \mid X_t, A_t).$$

An observation model gives:

$$Y_{t+1} \sim P(Y_{t+1} \mid X_{t+1}).$$

Nexus is therefore closer to a partially observable control problem than to a single-shot classifier.

A belief update takes the generic form:

$$b_{t+1}(x') \propto P(Y_{t+1} \mid x') \sum_x P(x' \mid x, A_t) b_t(x).$$

This matters because the meaning of a signal depends on context. The same EMG transient can correspond to different motor goals depending on gaze, world state, task, posture, and recent history. The system should therefore reason over joint evidence rather than treating one modality as a magical command channel.

5. The semantic operating system

5.1 Core objects

Nexus treats the following as first-class objects:

World
Context
Memory
Intent
Capability
Action
Policy
Compute
Trust
Device

This is the semantic kernel of the system.

World

A structured representation of relevant entities and relations:

$$W_t = (V_t, E_t, S_t)$$

where nodes can include people, places, objects, projects, tasks, documents, devices, events, concepts, goals, and capabilities.

Context

A relevance-weighted slice of world state conditioned on memory and task:

$$C_t = f(W_t, M_t, Q_t, G_t, H_t).$$

Memory

Typed persistent state rather than an undifferentiated chat transcript.

Intent

A probabilistic or structured representation of what the human is trying to accomplish.

Capability

An authorized operation that the environment can perform.

Action

A concrete execution request against a capability.

Policy

A mechanism for selecting or ranking candidate actions under uncertainty and constraints.

Compute

A resource abstraction spanning local processors, a pocket reservoir, nearby machines, wired hosts, and remote infrastructure.

Trust

The explicit set of permissions, provenance, confidence, reversibility, and risk constraints that separate “the system inferred this” from “the system is authorized to do this.”

5.2 Intent IR

A stable intermediate representation is essential if the architecture is to survive hardware and model changes.

A minimal object is:

```
@dataclass
class IntentIR:
    verb: str
    targets: list[str]
    parameters: dict
    context_refs: list[str]
    confidence: float
    urgency: float
    reversibility: str
    evidence: list[str]
```

Example:

```
{
  "verb": "compare",
  "targets": ["paper:A", "paper:B"],
  "parameters": {},
  "context_refs": ["project:photonics"],
  "confidence": 0.94,
  "urgency": 0.10,
  "reversibility": "high",
  "evidence": ["gaze", "voice", "context"]
}
```

The same Intent IR can be produced by voice, gaze plus a button, EMG, EEG evidence, keyboard input, or any combination.

This is the central decoupling:

$$\text{input modality} \neq \text{application semantics}$$

5.3 Action IR

Action IR is the inverse bridge. It turns a semantic action into executable device-independent structure.

Example:

`SELECT(entity=A)`

can lower to:

```
acquire gaze lock on A
render confirmation affordance
wait for valid confirmation
commit selection
```

and further to physical operations:

$$\{\theta_x(t), \theta_y(t), A(t), \phi(t), z_f(t)\}$$

for an optical system, or to device APIs for audio, robotics, desktop control, or communication.

5.4 Capability graph

The environment needs a description of what it can do:

Capability Graph

```
|
+-- communication
|   +-- send_message
|   +-- call
|   +-- share
|
+-- computation
|   +-- run_simulation
|   +-- train_model
|   +-- render
|
+-- world interaction
|   +-- inspect
|   +-- navigate
|   +-- locate
|   +-- manipulate
```

Providers can be native services, local programs, desktop applications, web services, external APIs, physical devices, or future Nexus hardware.

The system therefore inherits the existing digital world rather than rebuilding every capability from scratch.

5.5 Compatibility is a ladder, not a rewrite

Nexus should support:

Level	Description
L0	Existing application behaves normally
L1	Nexus can launch and route to it
L2	Specific capabilities become callable
L3	Application adopts Nexus semantic APIs
L4	Application is fully intent-native

This is strategically important. Platform transitions fail when they demand that the whole software ecosystem be rebuilt on day one.

6. Memory as a longitudinal personal substrate

A truly personal system cannot reset after every session.

6.1 Memory types

The minimum hierarchy is:

$$M = \{M_W, M_E, M_S, M_P, M_R, M_G\}$$

where:

- M_W : working memory;
- M_E : episodic memory;
- M_S : semantic memory;
- M_P : procedural memory;
- M_R : preference memory;
- M_G : goal and project state.

Every persistent memory item should carry provenance and policy metadata:

```
source
confidence
created_at
updated_at
scope
retention
user_control
sensitivity
```

The user should be able to inspect, edit, forget, export, or constrain memory.

6.2 Memory write policy

Persistent memory should not mean “store everything forever.”

A candidate score is:

$$S_{write} = \alpha N + \beta U + \gamma F + \delta I - \epsilon R_p$$

where:

- N : novelty;
- U : expected utility;
- F : future predictive value;
- I : explicit importance;

- R_p : privacy or sensitivity risk.

A memory write occurs only when the score and policy permit it.

This is not merely a storage optimization. It is an epistemic boundary: a system that remembers everything can become less useful than one that remembers selectively.

7. Human signal architecture

7.1 Multimodality is a design principle, not a feature list

Each modality answers a different question.

Channel	Main information
Gaze	attention direction / target
EMG	motor intention / subtle command
EEG	neural state or command evidence
Voice	semantic expression
IMU	head/body motion
World cameras	external scene state
Touch/button	explicit confirmation

The fused latent state is:

$$z_t = f_{\text{heta}}(EEG, EMG, Gaze, Voice, IMU, World, Memory).$$

The intent posterior is:

$$P(I_t | Y_t, C_t, M_t, Q_t).$$

The key idea is probabilistic evidence fusion. The modalities are not “voters” in a democracy. A modality can be wrong, missing, delayed, or corrupted. Reliability should be context-dependent.

A generic Bayesian formulation is:

$$P(I | Y, C) \propto P(Y | I, C)P(I | C).$$

7.2 Gaze

Gaze is valuable because it provides a spatial coordinate for attention. Its role should be modeled as more than cursor replacement.

Useful gaze signals include:

- point of regard;
- fixation stability;
- saccade dynamics;
- dwell time;
- vergence cues when available;
- blink state;
- pupil measurements where ethically and physically appropriate.

Gaze can disambiguate an otherwise weak neural or voice signal and can reduce the search space of intent inference.

7.3 EMG

Surface EMG is one of the most mature neural-adjacent channels in the Nexus stack because the signal is peripheral, closer to the muscles generating the action, and less ambiguous than unconstrained EEG for many discrete motor tasks.

External evidence supports the direction strongly. A 2025 Nature paper on a generic non-invasive neuromotor interface reported continuous navigation and gesture recognition using a wrist-worn sEMG system, including handwriting performance, demonstrating that the modality can support useful human-computer interaction at meaningful rates [EXT-07]. Meta has also publicly shipped an EMG wristband for subtle finger-driven interaction with its display glasses, providing product-scale evidence that the modality can enter everyday wearable use [EXT-10].

Meta Research’s emg2pose work provides a large-scale research resource for learning hand pose from sEMG, with recordings spanning many participants and held-out-user evaluation [EXT-08].

Nexus therefore does not need to prove that EMG exists. It needs to prove that the Nexus semantic stack can use it better as part of a persistent multimodal context model.

7.4 EEG

EEG provides richer access to neural state but arrives with major engineering costs:

- low signal amplitude;
- motion artifacts;
- variable contact impedance;
- nonstationarity;
- subject-specific structure;
- environmental electrical noise;
- long-term placement constraints.

A first research target is not “read arbitrary thought.” It is useful, bounded intent inference in known interaction domains.

The system should initially ask:

Does EEG add information beyond gaze + EMG + context?

rather than assume that it must.

The relevant benchmark is incremental information gain, not raw EEG classification accuracy.

7.5 Voice, vision, and IMU

Voice provides rich semantic evidence but is not always socially or acoustically appropriate. Vision provides world state but creates privacy and compute costs. IMU supplies motion reference and can improve the interpretation of gaze and biosignals.

The architecture should therefore treat every modality as optional, with graceful degradation when one fails.

8. Neuroengineering and bio-nano architecture

8.1 The hierarchy of biological interfaces

Nexus can be organized as a progression rather than a single commitment:

Level 0: conventional controls

Level 1: gaze + voice + IMU

Level 2: peripheral EMG

Level 3: non-invasive EEG + multimodal fusion

Level 4: high-density wearable neurophysiology

Level 5: research-grade minimally invasive interfaces

Level 6: future implantable / bidirectional systems

Each level adds possible bandwidth but also adds biological, regulatory, ethical, and engineering complexity.

The architecture therefore remains useful even if the deepest neural layers take decades to mature.

8.2 Electrode stack

For a future dry or semi-dry head-worn EEG path, the signal chain is:

```
skin
-> electrode interface
-> impedance/contact monitoring
-> low-noise biopotential AFE
-> ADC
-> digital filtering / artifact modeling
-> local representation encoder
-> secure transport
```

The TI ADS1299 family is an example of a commercially available low-noise multi-channel biopotential AFE class suitable for experimental architectures, not a claim that it is the final Nexus component [EXT-12].

Dry EEG literature indicates that dry electrodes can approach useful performance in some applications, while also emphasizing substantial variation in electrode designs and experimental methods [EXT-13]. That distinction matters: feasibility of a dry electrode modality is not the same as all-day glasses-scale reliability.

8.3 Candidate materials

The bio-nano materials program should evaluate families such as:

- conventional noble-metal electrodes;
- Ag/AgCl reference structures where appropriate;
- conductive polymer coatings including PEDOT:PSS;
- nanostructured metals;
- graphene and related carbon systems;
- MXene-like 2D conductive materials;
- soft and skin-conformal elastomer composites.

The selection function is multi-objective:

$$J_{bio} = \alpha SNR - \beta Z_{contact} - \gamma A_{motion} - \delta D_{drift} - \eta R_{skin} - \zeta C_{manufacturing}.$$

No material becomes useful because it is called “nano.” It becomes useful if the complete interface survives motion, sweat, repeated donning, long duration, manufacturing variance, and user-to-user variation while maintaining acceptable signal quality.

Current bioelectronic research supports flexible and conductive materials as promising platforms, while emphasizing long-term stability, mechanical mismatch, tissue or skin compatibility, and reliability as continuing challenges [EXT-15, EXT-16].

8.4 Future invasive path

The long-term research branch can include minimally invasive and implantable interfaces, but the product thesis should not depend on them.

Recent clinical and translational work demonstrates rapidly increasing channel density and useful closed-loop decoding in invasive neuroprostheses [EXT-17, EXT-18, EXT-19]. These results establish that high-bandwidth neural interfaces are scientifically real. They do not establish that an invasive system is consumer-ready, and they should not be folded into the Nexus v1 claim set.

The correct architecture is modular enough that a future high-bandwidth neural source can feed the same Intent IR without changing the entire operating environment.

9. Intelligence architecture

9.1 The assistant is a system process

The Nexus assistant is not intended to be another app opened by the user.

It is a controlled system process that can:

```
observe permitted context
retrieve memory
infer task state
predict candidate needs
resolve intent
select capabilities
propose actions
request confirmation when needed
execute authorized operations
record outcomes
learn under policy
```

The assistant does not own the hardware. The OS and safety supervisor do.

9.2 From prompt-response to persistent policy

The conventional assistant loop is:

$$prompt \rightarrow model \rightarrow answer.$$

Nexus aims for:

$$observe \rightarrow infer \rightarrow plan \rightarrow act \rightarrow observe \rightarrow learn.$$

The learned policy is approximately:

$$\pi_{\theta}(a \mid s, I, M).$$

An experience tuple is:

$$\tau_t = (s_t, a_t, r_t, s_{t+1}).$$

A practical product reward can combine success, interaction count, latency, interruption, correction, and trust:

$$r_t = \alpha R_{success} - \beta C_{interaction} - \gamma C_{latency} - \delta C_{error} - \eta C_{interrupt} - \kappa C_{trust}.$$

The system should never learn the absurd lesson that the easiest way to increase reward is to surprise the user.

9.3 Predictive interaction

The strongest user experience will often come from reducing the number of explicit actions.

A candidate prediction can be scored as:

$$U(a \mid c) = P(\textit{benefit} \mid c)V_b - P(\textit{interruption} \mid c)V_i.$$

Only actions whose utility exceeds a threshold should surface proactively, and the threshold should adapt to user preference without silently becoming lower because the system was rewarded for activity.

This turns “AI anticipation” into a measurable interaction policy.

10. Continual learning and post-training

10.1 Three adaptation levels

Nexus separates:

Level 1: memory adaptation

$$M_{t+1} = U_M(M_t, e_t).$$

Level 2: user-specific policy adaptation

$$\pi_u(a \mid s) = \pi_0(a \mid s) + \Delta\pi_\phi(a \mid s, M_t)$$

with bounded adaptation.

Level 3: global model improvement

Aggregate, privacy-filtered experience can contribute to slower global training or model selection.

The architecture deliberately separates the three because changing a user’s memory is not the same risk as changing a production model.

10.2 Safe post-training pipeline

experience log

- > privacy filter
- > episode construction
- > quality / label pass
- > counterfactual replay
- > simulator evaluation
- > candidate adapter
- > offline evaluation
- > canary deployment
- > monitoring
- > promotion / rollback

The system should preserve the previous policy and support rollback.

10.3 Why online RL must be bounded

Unbounded online reinforcement learning inside a consumer assistant creates a pathological incentive structure. The user becomes the environment, the reward function is partially observed, and mistakes can have real-world consequences.

Therefore:

$$\text{online adaptation} \subseteq \text{policy envelope}$$

A learned policy can choose within a trusted capability set. It cannot redefine permissions, bypass physical interlocks, install its own sensors, or rewrite safety constraints.

11. Recursive epistemic intelligence

Nexus uses a stronger concept of intelligence than prediction alone, but treats it explicitly as a research hypothesis.

11.1 Working definition

Recursive intelligence is an adaptive system's capacity to recursively improve the processes by which it forms representations, discovers abstractions, constructs and falsifies models, selects informative encounters and actions, compresses transferable knowledge, and reorganizes the computational mechanisms implementing those processes, under continuing interaction with an external reality.

The point is not self-modifying code. The point is the ability to change the mechanism by which the system learns what to learn.

11.2 The recursion ladder

A useful experimental taxonomy is:

$$R_1 \rightarrow R_2 \rightarrow R_3 \rightarrow R_4 \rightarrow R_5 \rightarrow R_6 \rightarrow R_7$$

with increasing scope:

1. state or memory adaptation;
2. model/error adaptation;
3. abstraction creation and revision;
4. representation revision;
5. question or goal refinement;
6. experiment / encounter selection;
7. epistemic or architectural self-revision.

This hierarchy is not assumed to be necessary. It is an ablation program.

11.3 Fixed search space versus changing the space

Ordinary optimization solves:

$$H^* = \arg \max_{H \in \mathcal{H}_t} J(H).$$

A more open system can change the hypothesis space:

$$\mathcal{H}_{t+1} = \Psi(\mathcal{H}_t, E_t).$$

An even stronger system may revise part of its evaluation structure:

$$J_{t+1} = \Psi_J(J_t, E_t).$$

That final capability must remain externally grounded. Otherwise an agent can become internally self-consistent while becoming less coupled to reality.

11.4 Falsification behavior

A critical test is whether the system deliberately chooses experiments that can distinguish competing models.

An epistemic experiment can be selected by:

$$a^* = \arg \max_a [\mathbb{E} (H(Q) - H(Q \mid Y_{t+1}, a)) - \lambda C(a)]$$

where $H(Q)$ is uncertainty over hypotheses.

The aim is not “more curiosity” in the abstract. The aim is efficient uncertainty reduction under physical and computational constraints.

12. World models and reality-facing intelligence

A Nexus intelligence system should maintain a compact predictive representation of the world rather than depend entirely on raw transcripts.

A latent representation is:

$$z_t = \Phi(Y_{0:t}, A_{0:t-1}).$$

A learned dynamics model predicts:

$$\hat{z}_{t+1} = f_\theta(z_t, a_t).$$

An ensemble can estimate epistemic uncertainty:

$$\mu(z_{t+1}) = \frac{1}{K} \sum_{k=1}^K f_k(z_t, a_t)$$

and dispersion can serve as one component of uncertainty.

The V-JEPA and world-model research ecosystem shows that predictive latent modeling is becoming a practical research direction for representation, prediction, and planning [EXT-20]. Nexus extends the idea into a persistent user-context environment, which remains its own research hypothesis.

The important distinction is:

$\text{memory} \neq \text{world model} \neq \text{policy}$

A system that confuses those three becomes hard to debug.

13. The compute reservoir

13.1 Why the reservoir exists

Trying to put the entire intelligence stack inside the glasses forces simultaneous optimization of:

$$\text{compute} + \text{battery} + \text{thermal} + \text{mass} + \text{RF} + \text{optics} + \text{wearability}.$$

The result is a set of conflicting constraints.

The reservoir changes the boundary:

glasses = sensing + optics + reflex

reservoir = cognition + memory + adaptation

remote = scale + training + long-horizon persistence

This is analogous to separating a sensor/actuator from the computational body of the system.

13.2 R0, R1, R2

The current architecture uses a family rather than a fixed silicon specification.

R0 - pocket reservoir

A phone-like or small custom compute appliance. Target envelope is a few watts of sustained compute for meaningful local intelligence, storage, radio, secure execution, and memory.

R1 - portable workstation reservoir

A larger mobile or mini-PC class unit with substantially more thermal headroom.

R2 - stationary reservoir

A desk, lab, or room-scale system suitable for high-throughput models, rendering, simulation, and local dataset work.

The exact power numbers remain design targets rather than commitments. The architecture should survive multiple silicon generations.

13.3 Compute placement as an optimization problem

Let a workload be mapped to node set N . Choose:

$$N^* = \arg \min_N C(N \mid L, B, P, E, R)$$

where:

- L : latency cost;
- B : bandwidth cost;
- P : privacy cost;
- E : energy cost;
- R : reliability cost.

A workload can move from glasses to reservoir to wired workstation to remote compute without changing its semantic identity.

13.4 Wireless and wired modes

Wireless is the default physical mode for everyday freedom.

Wired mode is not obsolete. It is the deterministic mode for:

- development;
- raw sensor capture;
- high-resolution optical debugging;
- simulation;
- long-duration heavy compute;
- fault analysis.

The logical stack remains identical.

The protocol classes are:

Class A - reflex

Clocking, optical state, safety status, high-priority gaze, timing synchronization.

Class B - interaction

Intent evidence, UI state, confirmations, control events.

Class C - intelligence

Model inference, memory retrieval, large data, training, rendering.

The critical rule is:

the inner physical control loop must not depend on an unreliable wireless round trip

13.5 Transport the minimum sufficient representation

Wireless transport should generally prefer compact, privacy-aware features over raw continuous biosignals when raw streams are not necessary.

For example:

`raw EEG -> local encoder -> uncertainty-aware latent -> reservoir`

rather than:

`raw EEG -> radio -> remote decoder -> action`

This reduces bandwidth, attack surface, power, and latency.

14. Power and thermal engineering

The reservoir exists because power is a systems variable.

Total consumption can be written:

$$P_{total} = P_{compute} + P_{memory} + P_{storage} + P_{RF} + P_{sensors} + P_{conversion} + P_{optics} + P_{misc}.$$

For the thermal state:

$$C_{\theta} \frac{dT}{dt} = P - hA(T - T_a).$$

At steady state:

$$\Delta T_{ss} \approx PR_{\theta}.$$

These equations are simple, but the engineering consequence is severe: a system that “fits” electrically for a benchmark can fail physically when operated continuously.

Nexus should therefore benchmark:

$$\frac{\text{useful task completion}}{\text{joules}}$$

rather than only TOPS, FPS, or model throughput.

15. Photonics: the physical interface as a computer

15.1 Optical architecture invariant

The long-term optical thesis is not “laser glasses.”

It is:

Make the delivered optical field computationally controllable and condition it on the state of the human, world, task, and device.

Possible implementations include:

- microLED + waveguide;
- laser/PIC + waveguide;
- MEMS laser scanning;
- optical phased arrays;
- diffractive or metasurface coupling;
- hybrid architectures.

The source technology is replaceable until experiments select it.

15.2 Field representation

A coherent optical field can be represented as:

$$E(x, y) = A(x, y)e^{i\phi(x, y)}.$$

The amplitude A controls magnitude; the phase ϕ controls interference structure.

A propagation operator can be expressed by angular-spectrum methods:

$$E_z = \mathcal{F}^{-1} \{ \mathcal{F}(E_0) H(f_x, f_y, z) \}.$$

These are standard wave-optics tools. They are not a Nexus invention. The research question is how to embed these physical transformations inside a real-time closed-loop system.

15.3 Waveguide physics

For a simple total-internal-reflection waveguide:

$$n_{core} > n_{clad}$$

and the critical angle satisfies:

$$\sin \theta_c = \frac{n_{clad}}{n_{core}}.$$

For a simplified grating coupling relation:

$$m\lambda = d(n_{eff} \sin \theta_i + \sin \theta_d).$$

Real waveguides involve dispersion, polarization, geometry, fabrication tolerances, angular response, pupil expansion, losses, and multiple coupled paths. The equations above define the first-principles skeleton, not the entire device.

15.4 Etendue and the unavoidable trade space

Optical throughput is bounded by etendue:

$$G = n^2 A \Omega.$$

You cannot arbitrarily increase field of view, eyebox, radiance, and compactness without paying for it somewhere in the optical system.

This is one of the fundamental reasons the display problem is hard.

15.5 External evidence for the direction

Recent optical research makes the integrated direction physically credible. A 2025 Nature Photonics paper demonstrated an ultra-thin mixed-reality display using a custom waveguide and AI-driven holography, reporting a 9 x 8 mm eyebox, 38 degree diagonal field of view, and diffraction-limited image characteristics while explicitly co-designing the optical hardware and computational algorithm [EXT-01].

A 2025 Nature paper demonstrated a flat-panel laser display built around a large-scale visible photonic integrated circuit and reported more than 80 percent volume reduction relative to a conventional LCoS architecture [EXT-02].

A 2025 Nature Nanotechnology result on single-layer waveguide displays demonstrated extreme thinness while highlighting the continuing trade-offs among chromatic control, efficiency, and uniformity [EXT-03].

These are not proof that Nexus is solved. They are proof that the physical research landscape already contains components of the proposed trajectory.

16. Adaptive optical control

16.1 State-dependent optical policy

Let:

$$x_t = \{G_t, W_t, C_t, Q_t, \Theta_t\}$$

contain gaze, world state, compute state, task requirements, and calibration state.

An optical policy proposes:

$$u_t = \pi_{opt}(x_t).$$

The optical system transforms this into an observed field:

$$y_t = F_{opt}(u_t, \Theta).$$

A controller uses the error:

$$e_t = y_t - y_t^*.$$

A simplified update is:

$$u_{t+1} = u_t + K(e_t).$$

A learned controller may replace or augment K , but only within bounded physical constraints.

16.2 Optical co-design with intelligence

The unusual Nexus hypothesis is:

AI does not merely choose what appears on the display; it helps decide how the optical system spends its physical resources.

Examples include:

- gaze-aware detail allocation;
- task-conditioned spatial bandwidth;
- context-aware field placement;
- confidence-dependent visual encoding;
- optical foveation;
- compute foveation;
- dynamic pupil targeting;
- adaptive calibration.

If this works, the system can optimize a different frontier than a static display.

16.3 Computational foveation

A conceptual rendering density map is:

$$R(x, y) = R_{min} + \alpha \exp \left(-\frac{\|[x, y] - g_t\|^2}{2\sigma^2} \right).$$

This concentrates computation around the gaze point.

When combined with optical resource allocation:

compute foveation + optical foveation

can become one jointly optimized pipeline.

This should be tested experimentally rather than assumed beneficial. Incorrect gaze, latency mismatch, or eye-tracking noise can make the result worse.

17. Optical see-through and the 3L interface

17.1 Current product architecture choice

The current Nexus product branch chooses optical see-through rather than camera passthrough as its primary architecture.

This is a product and human-factors decision, not a universal law of optics.

Optical see-through preserves direct world photons and makes the real world the primary optical substrate. Passthrough instead places a camera-plus-display reconstruction between eye and world.

The Nexus choice follows from a larger objective:

reality should remain the default visual substrate

while digital information is inserted only where useful.

17.2 Nexus 3L

The 3L interface is:

$$3L = Reality + Spatial\ Information + Intent$$

Layer 1 is the physical world.

Layer 2 is information anchored to objects, locations, paths, tasks, and relationships.

Layer 3 is intent-aware promotion: the system decides what information is relevant enough to enter conscious attention.

The goal is not to claim that resolution stops mattering. It remains a core physical parameter. The deeper claim is that pixels should be treated as a means to an interaction objective rather than as the entire interface definition.

18. Eye, depth, and perceptual constraints

A credible near-eye system must respect the human visual system.

Depth perception arises from multiple cues, including binocular disparity, accommodation, vergence, perspective, motion parallax, occlusion, and relative size. An optical system can reproduce some cues while failing to reproduce others.

This creates a systems problem:

$$\text{rendered depth} \neq \text{perceived depth automatically.}$$

A future system should therefore evaluate:

- focal behavior;
- vergence-accommodation interaction;
- stereo calibration;
- motion-to-photon latency;
- image stability;
- optical distortion;
- visual fatigue.

The right output metric is not “looks futuristic.” It is controlled human evidence.

19. Optical safety architecture

Safety is not a software popup.

If the system can expose an eye-facing optical source to hazardous conditions, safety must survive:

- eye-tracking loss;
- calibration failure;
- actuator freeze;
- software crash;
- corrupted command;
- power-control fault;
- network loss;
- sensor disagreement.

The system should therefore be layered:

```

learned optical policy
|
semantic supervisor
|

```

```

independent safety envelope
|
hardware limiter / source inhibit
|
physical optical path

```

For applicable coherent-source products, IEC TS 60825-20:2025 is directly relevant to face/eye-intentional laser exposure and explicitly addresses hazard analysis and failure mitigation for relevant headset/glasses classes [EXT-04]. ISO 15004-2:2024 provides a broader framework for optical radiation hazard protection in devices directing optical radiation into or at the eye, subject to product applicability [EXT-05].

The exact consumer product safety stack must be developed with the applicable standards and a formal hazard analysis; a mathematical barrier alone is not a substitute for certification.

The governing rule is:

learned systems may propose; deterministic safety systems decide

20. The bioelectronic safety hierarchy

The same separation applies to neurotechnology:

`observe != infer != suggest != execute != stimulate`

A sensor system can detect a signal without owning the right to act.

A decoder can infer intent without owning permission.

A planner can suggest an action without executing it.

A bidirectional stimulation system, if ever developed, adds a still higher risk tier.

This distinction should be encoded in the OS rather than left to developer discipline.

21. Security and neurodata sovereignty

21.1 Data classes

A possible privacy hierarchy is:

$$D = D_{raw} + D_{derived} + D_{intent} + D_{memory} + D_{cloud}.$$

Raw EEG, EMG, and eye imagery should be local by default wherever practical.

Derived intent should be exposed only to the capability that needs it.

Persistent memory should be inspectable and deletable.

The cloud should receive the minimum sufficient representation.

21.2 Threat model

Relevant threats include:

- neural data interception;
- replay attacks;
- decoder tampering;
- model poisoning;
- firmware compromise;
- malicious capabilities;
- unauthorized action execution;

- credential theft;
- privacy leakage from contextual memory;
- malicious cognitive manipulation;
- unsafe optical commands.

A baseline security chain is:

$$RootOfTrust \rightarrow SecureBoot \rightarrow KeyManagement \rightarrow EncryptedMemory \rightarrow CapabilityAccess \rightarrow AuditTrail.$$

21.3 Privacy as a physical architecture

Privacy is not simply a permissions page.

Local compute reduces the number of links across which sensitive data must travel.

Therefore:

$$\text{physical locality} \rightarrow \text{lower exposure surface}$$

This is one of the strongest arguments for the reservoir.

22. Human effort as a measurable systems variable

The human is part of the control loop.

Define:

$$C_{human} = w_v C_{visual} + w_c C_{cognitive} + w_s C_{social} + w_i C_{interaction}.$$

Then the product can optimize:

$$\max U_{task} - \lambda C_{human}$$

subject to physical and safety constraints.

Measurements should include:

- task completion time;
- error rate;
- explicit interaction count;
- attention switching;
- visual fatigue;
- cognitive workload;
- correction rate;
- trust calibration;
- social acceptability;
- long-duration comfort.

The system should not be allowed to claim “natural interaction” until these variables improve in controlled studies.

23. OS runtime and developer platform

The operating layer should expose APIs for:

Intent API

Context API

World API
 Memory API
 Capability API
 Action API
 Compute API
 Trust API
 Device API
 Optical API

The optical API should expose physical capabilities rather than emitter registers:

```
set_gaze_target()
set_virtual_depth()
set_focal_state()
set_eyebbox_target()
render_field()
set_luminance_bound()
set_temporal_profile()
request_optical_calibration()
query_optical_health()
enter_safe_state()
```

The compiler below the semantic API can change while the application remains stable:

$$\textit{Semantic Action} \rightarrow \textit{Intent/Action IR} \rightarrow \textit{Device Policy} \rightarrow \textit{Hardware Actuation}.$$

This is the same strategic lesson learned from successful hardware abstraction platforms: the durable interface is the software abstraction above the device, not the specific transistor arrangement below it.

24. The capability ecosystem

External integrations should enter through a capability graph rather than become the OS architecture themselves.

MCP-style tool protocols can be one integration route, but the Nexus architecture must remain protocol-agnostic.

The key distinction is:

$$\boxed{\textit{protocol} \neq \textit{architecture}}$$

The architecture defines what an intent, capability, action, memory, world state, and trust boundary mean. A protocol merely carries those semantics.

Developer tools should include:

```
nexus-cli
nexus-sim
nexus-debug
nexus-profiler
nexus-replay
nexus-inspector
nexus-benchmark
nexus-trace
```

The profiler should expose end-to-end traces:

$$T_{total} = T_{sensor} + T_{fusion} + T_{intent} + T_{plan} + T_{transport} + T_{render} + T_{optical}.$$

Without a trace model, latency becomes anecdotal and impossible to optimize.

25. The physical product architecture

The canonical physical system contains:

25.1 Glasses

Primary sensory and optical endpoint:

- optical-see-through display;
- world-facing camera;
- microphones;
- gaze sensing;
- IMU;
- candidate on-frame or near-frame EEG contacts;
- candidate EMG contacts for selected channels;
- local low-latency control electronics;
- high-speed wireless and wired connectivity.

25.2 Wrist band

Continuous surface EMG is the primary current candidate for subtle motor input.

This component deliberately exploits an emerging mature modality rather than attempting to invent every layer simultaneously.

25.3 Optional audio device

Earbuds or headphones can provide audio output and potentially future auxiliary sensing. The system behavior should not depend on one vendor's form factor.

25.4 Reservoir

Portable compute, secure memory, larger battery, radios, persistent storage, model runtime, personal memory, and learning infrastructure.

26. Distributed execution and graceful degradation

A robust Nexus system must keep functioning when parts disappear.

Define a dependency graph:

```
full intelligence
|
+-- reservoir
|   |
|   +-- wireless
|   +-- local memory
|
+-- glasses
    |
    +-- gaze
    +-- optics
    +-- voice
```

When the reservoir is unavailable, the glasses should retain a reduced local mode.

When EEG quality collapses, EMG + gaze + voice should continue.

When wireless fails, the physical optical loop should continue from locally available state within safe limits.

When a world model is uncertain, the system should ask for confirmation rather than hallucinate authority.
This is fault tolerance applied to human-machine interaction.

27. The core control law of Nexus

Across software, optics, and neurotechnology, one principle repeats:

$$\boxed{\text{infer} \rightarrow \text{bound} \rightarrow \text{act} \rightarrow \text{measure} \rightarrow \text{update}}$$

Inference alone is not execution.

Execution alone is not intelligence.

Learning alone is not truth.

The loop becomes scientifically meaningful only when action produces observations that can invalidate the model.

28. Benchmark suite

The benchmark suite should be end-to-end.

Let:

$$\mathcal{M} = \{A, T, B, E, R, C, V\}$$

where:

- A : intent accuracy;
- T : latency;
- B : explicit interaction burden;
- E : energy per task;
- R : safety / wrong-action risk;
- C : context continuity;
- V : visual quality.

A single scalar can summarize an engineering trade space for internal optimization, but public reports should preserve the vector because improvements in one dimension can hide regressions in another.

28.1 Intelligence benchmark

Compare:

1. prompt-response assistant;
2. context-aware assistant;
3. memory-aware agent;
4. predictive agent;
5. personalized policy;
6. full multimodal Nexus loop.

Measure:

- task time;
- interaction count;
- correction rate;
- interruption rate;
- energy;
- context recovery;
- transfer across tasks.

28.2 Neural benchmark

Report:

within-session
cross-session
cross-user
cross-environment
cross-device

Raw shuffled accuracy is insufficient.

EMG experiments should include cross-user and cross-session holdouts inspired by public research such as emg2pose [EXT-08]. EEG experiments should report dataset and session separation using reproducible public benchmarks such as PhysioNet motor movement/imagery and BCI Competition data [EXT-09, EXT-11].

28.3 Optical benchmark

Geometric

- FOV;
- eyebox;
- eye relief;
- distortion;
- pupil position;
- calibration drift.

Radiometric

- source power;
- optical efficiency;
- stray light;
- leakage;
- thermal load.

Photometric

- luminance;
- chromaticity;
- contrast;
- uniformity.

Wave-optic

- wavefront error;
- PSF;
- MTF;
- speckle;
- phase error.

Temporal

- source modulation latency;
- scan latency;
- motion-to-photon latency;
- jitter;
- synchronization error.

Human

- readability;
- comfort;
- fatigue;
- task success;
- accommodation/vergence behavior.

29. Decisive experiments

The thesis becomes real when it creates experiments that can destroy it.

Experiment 1 - Does context outperform sensor complexity?

Compare intent inference from:

$$EEG$$

versus:

$$EEG + EMG + Gaze$$

versus:

$$EEG + EMG + Gaze + Context + Memory.$$

Measure incremental information gain, false actions, and interaction burden.

A negative result would reduce the priority of deeper EEG investment.

Experiment 2 - Does prediction reduce total human effort?

Compare reactive and predictive policies.

Measure:

$$\Delta U_I, \Delta interruption, \Delta error.$$

Prediction is successful only if useful interaction falls without a compensating rise in annoyance or incorrect action.

Experiment 3 - Does personalization beat model scaling?

Hold model architecture approximately constant and compare:

- no memory;
- generic memory;
- user-specific memory;
- user-specific memory + bounded policy adaptation.

Measure long-horizon task performance and energy.

Experiment 4 - Can intelligent optics reduce resource cost?

Compare:

static rendering

vs

foveated rendering

vs
gaze + context aware rendering
vs
full adaptive optical control

Measure useful visual information per watt and task success per unit of compute.

Experiment 5 - Does the reservoir improve the wearable system frontier?

Compare the same glasses hardware under:

- all-local compute;
- split compute;
- reservoir-first compute.

Measure mass, skin temperature, battery life, latency, task success, and stability.

Experiment 6 - Can the intelligence discover and reject wrong models?

Use a controlled digital environment with hidden physical parameters.

Provide no explicit world model.

Measure whether the system:

1. forms a predictive model;
2. notices model error;
3. chooses a discriminating experiment;
4. updates the model;
5. transfers the resulting abstraction.

This connects Nexus to the broader recursive epistemic intelligence research program.

30. Validation ladder

The full program follows:

$$E_0 \rightarrow E_1 \rightarrow E_2 \rightarrow E_3 \rightarrow E_4 \rightarrow E_5 \rightarrow E_6$$

where:

- E0: first-principles equations and simulation;
- E1: component bench measurements;
- E2: hardware-in-loop and fault injection;
- E3: integrated head-worn prototype;
- E4: human controlled and longitudinal studies;
- E5: system validation, security, reproducibility, reliability;
- E6: production and lifecycle validation.

A claim may not legally or scientifically jump levels by wording.

31. Falsifiability and failure modes

31.1 The optical thesis fails if

The system cannot achieve an acceptable joint frontier of:

(FOV, eyebox, brightness, efficiency, color, MTF, latency, safety, mass, thermal).

A single record metric cannot rescue a system that fails the integrated frontier.

31.2 The neural thesis fails if

Neural channels do not provide useful incremental information after accounting for gaze, EMG, voice, context, and memory, at a power and comfort cost users will tolerate.

If this occurs, Nexus can remain intent-centric without deep neural input.

31.3 The intelligence thesis fails if

Predictive or personalized policies do not reduce interaction burden without increasing error or interruption, or if long-horizon memory does not improve task continuity.

31.4 The reservoir thesis fails if

The latency, thermal, battery, and connectivity penalties of split compute outweigh the advantages, or if commodity devices can provide the same abstraction with no meaningful loss of user experience.

31.5 The platform thesis fails if

Developers do not gain sufficient leverage from semantic capabilities over existing APIs, making the new abstraction expensive without producing a corresponding benefit.

These failure modes are not rhetorical disclaimers. They are the experimental escape hatches that keep the project honest.

32. Evidence from current science and engineering

The scientific foundation of Nexus is assembled from mature fields rather than one speculative breakthrough.

Domain	Established foundation	Nexus hypothesis
optics	Fourier/wave optics, waveguides, photonic integration	adaptive optical computer conditioned on context
photonics	integrated visible PICs and advanced waveguides	consumer-scale programmable field engine
biosignals	EEG/EMG acquisition and signal processing	multimodal intent layer with persistent personalization
EMG interaction	public research and product deployment	semantic motor channel inside a general OS
AI	world models, self-supervised representation learning, continual learning	persistent personal world/context model
RL	contextual policies and lifelong learning	bounded user-specific adaptation
OS	abstractions, IPC, drivers, security, compute schedulers	intent/capability/action as system primitives
security	secure boot, capability control, encryption	neurodata-aware action authority model
human factors	attention and workload measurement	conscious-effort optimization as a system objective

The project is therefore not built on the claim that any single technology is magic. Its novelty hypothesis sits in the integration.

33. The architecture moat

The moat should not be a list of hardware components.

A display can be copied.

A radio can be copied.

A compute puck can be copied.

An EMG band can be copied.

The harder-to-reconstruct asset is the longitudinal interaction substrate:

Intent IR + Capability Graph + Context Graph + Personal Memory + User Policy + Compute Fabric + Optical Interaction I

A competitor can reproduce a component and still fail to reproduce the accumulated state of the entire system.

This is an architectural moat, not a claim that the moat is already commercially defensible.

34. Open versus proprietary boundaries

The system should be open where openness increases network effects and scientific inspection.

Potentially open:

- Intent schemas;
- capability schemas;
- interoperability formats;
- benchmark definitions;
- developer documentation;
- data export formats;
- selected research tools.

Potentially proprietary:

- multimodal fusion models;
- personalization methods;
- context relevance engine;
- optical control policies;
- calibration algorithms;
- scheduler optimizations;
- hardware and firmware co-design;
- specialized silicon;
- manufacturing process knowledge.

The user's data should remain exportable.

The reason to stay should be capability and utility, not captivity.

35. Manufacturing and systems engineering

A device that works once in a lab is not a product.

The real stack is:

design → tolerance analysis → calibration → test → yield → serviceability

Optical systems are especially sensitive to manufacturing variance. Neural interfaces are sensitive to contact geometry and user variation. Wireless systems are sensitive to body shadowing and antenna placement. Thermal design is

sensitive to skin proximity and frame packaging.

Therefore every future Nexus subsystem needs a tolerance budget.

For a generic system parameter x :

$$\sigma_{system}^2 = \sum_i \left(\frac{\partial f}{\partial x_i} \right)^2 \sigma_i^2$$

provides a first-order propagation of manufacturing uncertainty.

Calibration should be treated as part of architecture, not a factory afterthought.

36. The product evolution path

The long-term architecture should be developed in stages so that each stage can stand on its own.

Stage A - semantic Nexus on commodity hardware

Build Intent IR, Capability Graph, memory, compute fabric, trace tooling, developer APIs, and a context-aware assistant over existing devices.

Stage B - multimodal wearable interaction

Integrate gaze, IMU, voice, EMG, and spatial presentation. Measure the interaction advantage.

Stage C - reservoir-native execution

Move persistent memory and heavier inference to R0. Demonstrate meaningful local privacy and lower wearable burden.

Stage D - advanced optical co-design

Integrate a research optical engine whose resource allocation is conditioned on task and human state.

Stage E - neural expansion

Add validated EEG pathways where experiments demonstrate incremental information gain.

Stage F - longitudinal intelligence

Demonstrate weeks-to-months personalization, rollback, memory governance, and measurable improvements in task continuity.

Stage G - advanced neurotechnology research

Only if scientifically and ethically justified, investigate higher-bandwidth minimally invasive or implantable interfaces.

The staged strategy matters because it prevents the project from becoming dependent on the hardest unsolved subsystem.

37. The research operating model

Nexus should be run as an evidence machine.

For every subsystem, maintain:

```

claim
-> equation / mechanism
-> baseline
-> experiment
-> measurement
-> failure criterion
-> artifact
-> decision

```

A research log should preserve not only successes but rejected designs.

A failed optical architecture can be scientifically valuable if it reveals which constraint dominated.

A neural model that fails across sessions can be more informative than a high within-session score if it exposes the stability bottleneck.

A policy that predicts the wrong action once can become a valuable test case if the trace reveals why the system over-weighted a misleading context.

This is how the project keeps its hunger while remaining scientific.

38. The engineering creed

The program can be compressed into ten rules.

1. **Build the smallest artifact that collapses the biggest uncertainty.**
2. **Never promote simulation to reality by adjectives.**
3. **Use the simplest model that can be falsified.**
4. **Treat every learned component as untrusted until bounded.**
5. **Keep physical safety outside the learned policy.**
6. **Measure the human, not only the machine.**
7. **Prefer architectures that degrade gracefully.**
8. **Make memory inspectable and reversible.**
9. **Change the abstraction only when the evidence demands it.**
10. **Turn every strong belief into an experiment.**

39. Master system optimization

All Nexus subsystems can be expressed as a constrained optimization:

$$\begin{array}{ll}
 \max_{\pi, u, N, M, O} & U_{human}(task, context) \\
 \text{s.t.} & S_{optical} \leq S_{max} \\
 & T_{skin} \leq T_{max} \\
 & P_{battery} \leq P_{max} \\
 & \epsilon_{sync} \leq \epsilon_{sync, max} \\
 & \epsilon_{gaze} \leq \epsilon_{gaze, max} \\
 & R_{action} \leq R_{max} \\
 & C_{neural} \geq C_{min} \\
 & Q_{image} \geq Q_{min} \\
 & M_{device} \leq M_{max} \\
 & \text{Privacy, permission, trust constraints hold.}
 \end{array}$$

This formulation is intentionally broader than a display optimization or AI benchmark. It asks whether the integrated

system is useful under real physical and human constraints.

40. The single deepest bottleneck

The hardest problem is not one component.

It is synchronization of heterogeneous timescales and confidence levels.

An EEG decoder may operate over hundreds of milliseconds.

A gaze estimator may update faster.

A world model may update at a different rate.

A wireless link may introduce variable latency.

An optical actuator can require deterministic control.

The user expects one coherent experience.

Therefore the true systems problem is:

deterministic coordination across uncertain, asynchronous subsystems


This is why Nexus must be treated as a systems engineering project rather than an accumulation of fashionable components.

41. Why the architecture can still be radical without pretending everything is solved

The strongest version of the Nexus claim is not that every ingredient is novel.

Most ingredients are not.

The radical proposition is that the ingredients can be composed around a different system boundary:



This shifts computing from a device-centered abstraction toward an environment-centered one.

It changes the question from:

“How do I operate this computer?”

to:

“What capability do I need, and what is the lowest-friction safe path to it?”

That is a sufficiently deep architectural change to justify a full research program even before the consumer hardware is complete.

42. Public research agenda

The next generation of Nexus research should produce artifacts, not only documents.

Software

- stable Intent IR implementation;
- Action IR compiler;
- Capability Graph runtime;
- Context Graph;
- typed memory fabric;
- compute scheduler;
- end-to-end trace system;
- safe policy envelope;
- replay and benchmark harness.

Intelligence

- multimodal encoder;
- user-state estimator;
- world model;
- memory relevance model;
- contextual policy learner;
- counterfactual evaluator;
- epistemic experiment selector.

Neurotechnology

- synchronized EEG/EMG/gaze capture;
- artifact models;
- cross-session benchmarks;
- cross-user adaptation;
- electrode prototypes;
- secure neural transport.

Photonics

- wave-optics simulator;
- parameterized waveguide model;
- source/steering co-simulation;
- optical control policy;
- calibration model;
- hardware-in-loop photodiode / eye model;
- safe source-control architecture.

Systems

- R0 reservoir prototype;
- thermal instrumentation;
- wired/wireless transport;
- secure boot chain;

- deterministic reflex loop;
- power-aware scheduler.

43. What Nexus is, finally

Nexus is not fundamentally a pair of glasses.

It is not fundamentally an EEG device.

It is not fundamentally an AI assistant.

It is not fundamentally a compute puck.

It is not fundamentally a photonic display.

It is a **human-compute operating environment** whose physical manifestations are designed together.

The software makes human intent, context, memory, capability, and computation first-class.

The reservoir makes persistent intelligence physically practical.

The neural and physiological layer expands the input bandwidth where it proves useful.

The photonic layer turns output into spatially coherent, adaptive information.

The control layer keeps fast physical behavior deterministic.

The security layer ensures inference does not equal authority.

The learning layer allows the environment to become more useful to one person over time.

The research program allows every one of those statements to be tested rather than merely believed.

44. Final compression

The complete architecture can be reduced to one dynamical equation set:

$$\begin{aligned}
 H_{t+1} &= F_H(H_t, A_t, W_t) \\
 W_{t+1} &= F_W(W_t, A_t) \\
 M_{t+1} &= F_M(M_t, H_t, W_t, A_t) \\
 I_t &= F_I(H_t, W_t, M_t, Q_t) \\
 A_t &= F_A(I_t, M_t, C_t, T_t) \\
 O_t &= F_O(A_t, G_t, \Theta_t) \\
 G_{t+1} &= F_G(G_t, O_t) \\
 C_{t+1} &= F_C(C_t, \text{workload}, \text{network}, \text{thermal})
 \end{aligned}$$

with safety and authority constraints:

$$A_t \in \mathcal{A}_{\text{authorized}}(T_t, \text{Trust}_t),$$

and physical feasibility:

$$O_t \in \mathcal{O}_{\text{safe}}.$$

The system is intelligent only insofar as the loop produces useful, robust, transferable behavior and remains coupled to reality.

The deepest thesis is therefore:

NEXUS = persistent context + multimodal intent + capability execution + adaptive computation + program

And the deepest engineering principle is:

\$\$ **Do not build a machine that asks the human to translate reality into software. Build a machine that learns**

The ambition is large. The standard must remain unforgiving.

The goal is not to win an argument about what the future will look like.

The goal is to make the future submit to measurement.

45. A researcher’s reading of Nexus

Nexus has four different faces, and the thesis should remain legible from all four.

Reader	First question	Nexus answer
Scientist	What is the falsifiable hypothesis?	The integrated human-state / world-state / compute / optical loop produces measurable system-level gains under physical constraints.
Engineer	What are the interfaces and failure boundaries?	Intent IR, Action IR, Capability Graph, Device Graph, transport classes, timebase, deterministic control, safety envelope.
Builder	What can be built now?	Software semantic runtime, reservoir prototype, multimodal capture, optical simulation, hardware-in-loop instrumentation.
Investor / fellowship reviewer	Why is this more than a feature?	The proposed value sits at the architecture boundary: a persistent personal compute environment with a new semantic interface and a co-designed physical endpoint.

The thesis deliberately avoids saying that the final integrated device already exists. What exists is a coherent engineering program with enough software and simulation structure to start collapsing the remaining uncertainties experimentally.

46. Cognitive runtime, attention, and externalized working memory

The original screen-centric computer makes the human maintain a large amount of state:

where was I?
which window?
which document?
which object?
which step?
what did I already tell the system?
what was the previous result?
what should happen next?

A persistent Nexus environment should absorb part of this burden.

The cognitive runtime therefore maintains:

$$K_t = (Task_t, Context_t, Memory_t, Attention_t, Goal_t).$$

The system should estimate the relevance of an event to this state rather than surface every event equally.

A simple attention allocation model is:

$$A_i = w_r R_i + w_u U_i + w_t T_i + w_c C_i - w_d D_i$$

where relevance, expected utility, urgency, contextual fit, and distraction cost are combined into a presentation priority.

This does not mean the machine owns the user's attention. It means the system treats attention as a scarce resource and optimizes how much it consumes.

46.1 Externalized working memory

A user should be able to move through a long task without manually reconstructing the state of the computer.

For a scientific workflow, for example:

```
problem definition
-> literature state
-> active derivation
-> simulation parameters
-> experimental result
-> unresolved anomaly
-> next experiment
```

Nexus can persist the structure rather than only the textual transcript.

The long-term research question is whether this externalized state reduces working-memory burden without reducing understanding.

That is testable with controlled experiments comparing:

ordinary desktop workflow

versus

Nexus context/memory workflow.

46.2 Attention must remain user-steerable

A predictive assistant that reduces clicks but constantly interrupts is not a success.

Therefore every proactive behavior should expose:

- confidence;
- reason for surfacing;
- reversibility;
- suppression controls;
- learned preference state.

The user should be able to make the system less proactive without losing core access to the rest of the environment.

47. Applications as experimental environments

The first applications should be chosen not because they sound futuristic, but because they make the architectural advantage measurable.

47.1 Scientific research

A scientist looks at a paper, a simulation, a microscope image, or a physical apparatus. Nexus maintains project context and can place a relevant computation or explanation spatially near the object under attention.

A possible workflow is:

```
look at equation
-> identify symbol / object
-> retrieve project context
-> execute symbolic or numerical capability
-> render result in place
-> record outcome
```

This directly tests whether spatial context and memory reduce application switching.

47.2 CAD and engineering

A designer can inspect a component and request:

```
check tolerance
show failure modes
compare material
run thermal simulation
```

The capability graph maps those requests onto local software, workstation compute, or remote solvers.

This becomes a strong testbed because the underlying digital artifacts are structured and the requested actions have measurable outcomes.

47.3 Education

Nexus can attempt to make the environment itself the object of instruction:

```
look at apparatus
-> identify component
-> visualize field / force / trajectory
-> manipulate parameters
-> observe model consequence
```

The research question becomes whether spatially anchored explanations improve learning or merely add visual novelty.

47.4 Communication

Communication is a useful test for contextual memory and prediction:

```
person
-> identity/context
-> recent interaction
-> relevant task
-> appropriate communication action
```

The system must remain conservative because incorrect context assumptions can be socially costly.

47.5 Industrial and field operations

Field workers often operate with hands occupied, limited attention, intermittent connectivity, and expensive mistakes.

This is an especially clean test of the reservoir architecture because:

hands-free interaction + local compute + context + degraded-network operation

are all simultaneously relevant.

48. What is deliberately not the novelty claim

Nexus should never claim novelty merely because it uses words that are individually fashionable.

Not enough by itself	Why
AI assistant	widely established category
smart glasses	existing product class
EEG	mature research modality
EMG wristband	emerging commercial modality
waveguide display	established AR research direction
laser projector	mature optical technology
photonic integrated circuit	established research / industry area
reinforcement learning	mature machine-learning field
persistent memory	active agent-systems research
cloud + edge compute	established distributed-computing pattern

The actual research claim is the composition:

persistent personal state + multimodal intent + capability-centric OS + distributed compute + adaptive physical presentation +

The burden of proof is therefore system-level.

49. Current Nexus branch decisions versus architectural invariants

Some design choices are currently strong commitments because they simplify the engineering path. They are not laws of nature.

Current product branch

1. Optical see-through is the primary display architecture.
2. Heavy intelligence is split outward from the glasses.
3. A pocket reservoir is the preferred center of personal compute.
4. A wrist EMG device is the practical near-term continuous motor channel.
5. EEG is a research expansion channel, not a prerequisite for the first useful product.
6. Glass waveguides are a practical initial optical material class, with higher-index material research kept on the roadmap.
7. Wireless is the consumer path; wired is the deterministic research/developer path.

Architectural invariants

The actual invariants are stronger than those component choices:

intent is independent of input modality	
capability is independent of device identity	
compute is movable across the fabric	
memory is persistent but governed	
learned policy cannot bypass authority or safety	
physical control remains deterministic	optical output is part of the computational loop

If a different technology achieves these invariants more effectively, the component architecture should change.

50. Master evidence matrix

Subsystem	Current Nexus state	External basis	Decisive next proof
semantic OS	N1 implemented architecture	OS / distributed-systems principles	cross-device intent execution
memory fabric	N1 design + prototype direction	lifelong-agent memory literature	longitudinal reduction in repeated user effort
compute fabric	N1 design	distributed computing + edge inference	measurable energy/latency benefit versus single-node baseline
EMG	integration architecture	product + published sEMG evidence	cross-user, cross-session semantic intent benchmark
EEG	research architecture	public EEG datasets + wearable BCI literature	incremental information gain after multimodal fusion
optical engine	N2 simulation + architecture	wave optics + modern waveguide/PIC research	calibrated hardware-in-loop optical control
intelligent optics	H1	gaze-aware rendering and adaptive holography literature	useful information per watt / task improvement
reservoir	N2 prototype architecture	mobile/edge accelerator ecosystem	all-day task-energy and thermal validation
recursive intelligence	H1 theory + simulator direction	world-model / lifelong-learning research	unknown-world epistemic reconstruction benchmark
safety	N1 architecture	laser safety + secure-systems principles	independent fault-injection validation

This matrix is the project's public contract with reality: the unknown rows are supposed to remain unknown until experiments move them.

Appendix N. Formal state machine for a Nexus action

A useful action state machine is:

```

OBSERVE
|
v
ESTIMATE
|
v

```

```

INTENT_CANDIDATE
|
+---- low confidence ----> ASK / WAIT
|
v
CAPABILITY_RESOLUTION
|
v
RESOURCE_ASSIGNMENT
|
v
TRUST_CHECK
|
+---- denied -----> CANCEL
|
v
EXECUTE
|
v
OBSERVE_OUTCOME
|
+---- failure -----> DIAGNOSE / RECOVER
|
v
MEMORY_UPDATE
|
v
POLICY_EVALUATION

```

The learned model participates primarily in ESTIMATE, CANDIDATE, RESOURCE, and POLICY stages. Hard permission, safety, and physical control remain outside the adaptive core.

Appendix O. Failure taxonomy

Nexus failures should be labeled by abstraction level.

F0 - sensor failure

Signal missing, corrupted, saturated, or stale.

F1 - estimation failure

The signal is valid but the inferred state is wrong.

F2 - context failure

The current world/task context is incomplete or stale.

F3 - model failure

The predictive model fails under an unmodeled condition.

F4 - policy failure

The selected action is poor even though state estimation was reasonable.

F5 - capability failure

The requested operation is unavailable or misconfigured.

F6 - transport failure

The logical operation is correct but the physical link is late or absent.

F7 - actuation failure

The physical device does not realize the commanded state.

F8 - safety failure

The system permits an unsafe state to persist.

The purpose of this taxonomy is to avoid the unhelpful diagnosis “the AI failed.” A complex system requires localized failure attribution.

Appendix P. Resource-aware inference policy

A model-selection policy can use:

$$J(m, q) = \alpha L(m, q) + \beta E(m, q) + \gamma B(m, q) + \delta R(m, q) - \eta Q(m, q)$$

where latency, energy, bandwidth, risk, and expected answer quality are normalized according to task requirements.

For a safety-critical or privacy-critical request, the weights change.

This is the principle of **conditional intelligence**:

use the minimum computational machinery that reliably solves the task

The architecture should not confuse intelligence with model size.

Appendix Q. Research log schema

Each experiment should record:

```

experiment_id
hypothesis
baseline
hardware_revision
software_revision
dataset
train/test split
random seed
environment
sensor configuration
model configuration
power profile
thermal profile
latency profile
result
confidence interval
failure cases
interpretation

```

next decision
artifact links

This makes the project auditable and reproducible.

Appendix R. Final public disclosure

Nexus is an ambitious engineering hypothesis spanning several fields whose individual components are at different maturity levels. Some pieces are established; some are demonstrated by external research; some are implemented as software or simulations; some remain entirely open.

The most important intellectual property of the project is therefore not the assertion that every piece is already solved. It is the systems hypothesis that these pieces can be organized into a new human-compute boundary and that the resulting advantage can be measured.

A successful Nexus prototype will not be the one that contains the most technologies.

It will be the one that proves, quantitatively, that the human needs to do less interface work while the computer does more of the coordination - without sacrificing agency, safety, privacy, or physical coherence.

That is the standard by which the thesis should be judged.

Appendix A. Minimal Intent IR schema

```
{
  "intent_id": "uuid",
  "verb": "select",
  "targets": ["entity:42"],
  "parameters": {},
  "context_refs": ["task:current"],
  "confidence": 0.91,
  "urgency": 0.20,
  "reversibility": "high",
  "evidence": [
    {"source": "gaze", "weight": 0.54},
    {"source": "voice", "weight": 0.31},
    {"source": "context", "weight": 0.15}
  ],
  "timestamp": "...",
  "provenance": "sensor-fusion-vX"
}
```

Appendix B. Minimal Action IR schema

```
{
  "action_id": "uuid",
  "capability": "document.compare",
  "arguments": {"left": "paper:A", "right": "paper:B"},
  "authorization": "policy-token",
  "risk": "low",
  "reversibility": "high",
  "resource_hint": "reservoir",
  "presentation": {
    "mode": "spatial",
    "priority": "normal"
  }
}
```

```
}
}
```

Appendix C. Protocol and compute classes

Class	Purpose	Typical payload	Timing character
A	reflex	timing, gaze, safety, optical state	deterministic / highest priority
B	interaction	intent evidence, confirmations	low latency
C	intelligence	models, memory, heavy data	throughput-oriented

The transport can be wireless or wired without changing the semantic ABI.

Appendix D. Neurotechnology benchmark protocol

The minimum reproducible protocol should separate:

1. training users from held-out users;
2. training sessions from test sessions;
3. development devices from evaluation devices;
4. calibration time from steady-state performance;
5. clean data from realistic motion.

Report:

Accuracy, FPR, FNR, latency, calibration time, energy, robustness, user retention.

The most important result is not an isolated accuracy number. It is whether useful semantic performance survives the conditions of actual wear.

Appendix E. Optical parameterization

A candidate simulator should expose:

```
wavelength
source spectrum
beam divergence
phase map
waveguide index
waveguide thickness
grating period
grating depth
coupling efficiency
FOV
eyebow
eye relief
pupil diameter
eye position
head pose
luminance
speckle
```

surface roughness
manufacturing tolerance

The system should generate both optical outputs and sensitivity information:

$$\frac{\partial Q}{\partial p_i}$$

for each design parameter p_i .

Appendix F. Reservoir thermal model

A minimal instrumented model:

$$P(t) = P_{compute} + P_{RF} + P_{memory} + P_{sensor} + P_{conversion}$$

$$C_\theta \frac{dT}{dt} = P(t) - hA(T - T_a).$$

Measure:

- CPU package temperature;
- skin-adjacent temperature;
- reservoir enclosure temperature;
- optical source temperature;
- battery temperature;
- radio duty cycle;
- workload occupancy.

Then optimize task utility per joule rather than compute throughput alone.

Appendix G. Current public evidence sources

The following references support the external scientific and engineering evidence cited in this thesis. They are intentionally separated from the Nexus-internal corpus.

[EXT-01] Choi et al., “Synthetic aperture waveguide holography for compact mixed-reality displays with large etendue,” Nature Photonics (2025). <https://www.nature.com/articles/s41566-025-01718-w>

[EXT-02] “Flat-panel laser displays through large-scale photonic integrated circuits,” Nature (2025). <https://www.nature.com/articles/s41566-025-09107-7>

[EXT-03] “Single-layer waveguide displays using achromatic metagratings,” Nature Nanotechnology (2025).

[EXT-04] IEC TS 60825-20:2025, Safety of laser products - Part 20, eye/face intentional exposure applications including relevant headset/glasses classes.

[EXT-05] ISO 15004-2:2024, Ophthalmic instruments - Fundamental requirements and test methods - Part 2: Optical radiation hazard protection.

[EXT-06] ISO 10993-1:2025, Biological evaluation of medical devices - Part 1.

[EXT-07] “A generic non-invasive neuromotor interface for human-computer interaction,” Nature (2025).

[EXT-08] Meta Research, emg2pose dataset and research repository. <https://github.com/facebookresearch/emg2pose>

[EXT-09] PhysioNet EEG Motor Movement/Imagery Dataset. <https://physionet.org/content/eegmmidb/1.0.0/>

[EXT-10] Meta, Neural Band product/research materials, 2025-2026.

[EXT-11] BCI Competition IV. <https://www.bbci.de/competition/iv/>

- [EXT-12] Texas Instruments ADS1299 product documentation.
- [EXT-13] Review literature on dry EEG electrodes and wearable EEG acquisition.
- [EXT-14] Flexible brain electronic sensors and wearable BCI review literature, Nature family journals (2025).
- [EXT-15] PEDOT:PSS-based bioelectronics for brain monitoring and modulation, Nature (2025).
- [EXT-16] Nature Reviews Electrical Engineering / Bioengineering literature on skin-conformal electronics and low-dimensional bioelectronic materials (2025-2026).
- [EXT-17] “Minimally invasive implantation of scalable high-density cortical microelectrode arrays,” Nature Biomedical Engineering (2026).
- [EXT-18] “An instantaneous voice-synthesis neuroprosthesis,” Nature (2025).
- [EXT-19] “Long-term independent use of an intracortical BCI,” Nature Medicine (2026).
- [EXT-20] Meta AI, V-JEPA 2 technical materials (2025).
- [EXT-21] Memory OS of AI Agent, EMNLP (2025).
- [EXT-22] Lifelong reinforcement learning knowledge preservation literature, Nature Machine Intelligence (2025).
- [EXT-23] USB4 Specification Version 2.0 / current USB4 family documentation.
- [EXT-24] Qualcomm Snapdragon AR2 Gen 1 platform architecture and vendor reference claims.

Appendix H. Nexus internal source corpus integrated into this edition


This master thesis synthesizes the following internal Nexus documents and research streams available in the project corpus:

- NEXUS_OS_README.md
- 04_Nexus_Computing_Reservoir_Whitepaper.md
- 06_Nexus_Computing_Reservoir_Engineering_Blueprint.md
- 03_Nexus_Full_Stack_Human_Compute_Photonics_Architecture.md
- 07_Nexus_All_In_One_Product_System_Architecture.md
- MASTER_TECHNICAL_RND_WHITEPAPER.md
- Project_NEXUS_Research_Paper.md
- Project_NEXUS_Research_Log.md
- recursive_epistemic_intelligence_paper.md
- NEXUS_Thiel_Fellowship_Interview_Dossier.pdf
- NEXUS - Master Execution Plan - Network, Engineering, Research, Manufacturing, Safety, Product, Platform & Global Scale.md
- related Nexus product-constitution and engineering notes in the project corpus.

The purpose of the master edition is not to preserve every source sentence. It preserves the substantive architecture, equations, engineering arguments, research questions, evidence boundaries, and execution logic in one public object.

Appendix I. Master architecture figure


NEXUS distributed human-compute operating environment



Measured outcome returns to state estimation; physical safety remains independent of learned policy.

Figure 1: Nexus distributed human-compute architecture

Appendix J. Three-timescale figure




The layers exchange targets and measurements, but each retains its own timing and authority boundary.

Figure 2: Three-timescale control architecture

Appendix K. Photonic architecture figure

Adaptive photonic near-eye engine: research architecture




Invariant: the delivered optical field is computationally controllable.
The source, steering and coupling technology remain experimentally selectable.

Figure 3: Adaptive photonic stack

Appendix L. Learning loop figure

Personal intelligence: bounded learning and memory loop




Online adaptation is reversible, bounded, auditable and user-controlled.

Figure 4: Personal learning loop

Appendix M. Evidence ladder figure

Nexus evidence ladder



**A design is not a demonstration. A simulation is not a measurement.
A measurement is not product certification.**

Every public claim should point to the highest evidence level that actually exists.

Figure 5: Evidence ladder